

```
In [1]: import numpy as np
import pandas as pd
```

```
In [2]: crop = pd.read_csv(r"C:\Users\Utkarsh\Desktop\Juggad_hacks\Crop_recommendation.csv")
crop.head()
```

```
Out[2]:
```

	N	P	K	temperature	humidity	ph	rainfall	label
0	90	42	43	20.879744	82.002744	6.502985	202.935536	rice
1	85	58	41	21.770462	80.319644	7.038096	226.655537	rice
2	60	55	44	23.004459	82.320763	7.840207	263.964248	rice
3	74	35	40	26.491096	80.158363	6.980401	242.864034	rice
4	78	42	42	20.130175	81.604873	7.628473	262.717340	rice

```
In [3]: crop.shape
```

```
Out[3]: (2200, 8)
```

```
In [4]: crop.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2200 entries, 0 to 2199
Data columns (total 8 columns):
 #   Column          Non-Null Count  Dtype
---  -
 0   N                2200 non-null   int64
 1   P                2200 non-null   int64
 2   K                2200 non-null   int64
 3   temperature      2200 non-null   float64
 4   humidity         2200 non-null   float64
 5   ph               2200 non-null   float64
 6   rainfall         2200 non-null   float64
 7   label           2200 non-null   object
dtypes: float64(4), int64(3), object(1)
memory usage: 137.6+ KB
```

```
In [5]: crop.isnull().sum()
```

```
Out[5]:
```

N	0
P	0
K	0
temperature	0
humidity	0
ph	0
rainfall	0
label	0

dtype: int64

```
In [6]: crop.duplicated().sum()
```

```
Out[6]: 0
```

```
In [7]: crop.describe()
```

Out[7]:

	N	P	K	temperature	humidity	ph	rainfall
count	2200.000000	2200.000000	2200.000000	2200.000000	2200.000000	2200.000000	2200.000000
mean	50.551818	53.362727	48.149091	25.616244	71.481779	6.469480	103.463655
std	36.917334	32.985883	50.647931	5.063749	22.263812	0.773938	54.958389
min	0.000000	5.000000	5.000000	8.825675	14.258040	3.504752	20.211267
25%	21.000000	28.000000	20.000000	22.769375	60.261953	5.971693	64.551686
50%	37.000000	51.000000	32.000000	25.598693	80.473146	6.425045	94.867624
75%	84.250000	68.000000	49.000000	28.561654	89.948771	6.923643	124.267508
max	140.000000	145.000000	205.000000	43.675493	99.981876	9.935091	298.560117

In [8]:

```
corr = crop.corr()  
corr
```

Out[8]:


	N	P	K	temperature	humidity	ph	rainfall
N	1.000000	-0.231460	-0.140512	0.026504	0.190688	0.096683	0.059020
P	-0.231460	1.000000	0.736232	-0.127541	-0.118734	-0.138019	-0.063839
K	-0.140512	0.736232	1.000000	-0.160387	0.190859	-0.169503	-0.053461
temperature	0.026504	-0.127541	-0.160387	1.000000	0.205320	-0.017795	-0.030084
humidity	0.190688	-0.118734	0.190859	0.205320	1.000000	-0.008483	0.094423
ph	0.096683	-0.138019	-0.169503	-0.017795	-0.008483	1.000000	-0.109069
rainfall	0.059020	-0.063839	-0.053461	-0.030084	0.094423	-0.109069	1.000000

In [9]:

```
import seaborn as sns  
sns.heatmap(corr,annot=True,cbar=True, cmap='coolwarm')
```

C:\Users\Utkarsh\anaconda3\lib\site-packages\scipy__init__.py:155: UserWarning: A NumPy version >=1.18.5 and <1.25.0 is required for this version of SciPy (detected version 1.26.4
warnings.warn(f"A NumPy version >={np_minversion} and <{np_maxversion}")
<AxesSubplot:>

Out[9]:




```
In [10]: crop['label'].value_counts()
```

```
Out[10]: rice          100
maize          100
jute           100
cotton         100
coconut        100
papaya         100
orange         100
apple          100
muskmelon      100
watermelon     100
grapes         100
mango          100
banana         100
pomegranate    100
lentil         100
blackgram      100
mungbean       100
mothbeans      100
pigeonpeas     100
kidneybeans    100
chickpea       100
coffee        100
Name: label, dtype: int64
```

```
In [11]: import matplotlib.pyplot as plt
sns.distplot(crop['N'])
plt.show()
```

C:\Users\Utkarsh\anaconda3\lib\site-packages\seaborn\distributions.py:2619: FutureWarning: `distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).
warnings.warn(msg, FutureWarning)



```
In [12]: crop_dict = {
    'rice': 1,
    'maize': 2,
    'jute': 3,
    'cotton': 4,
    'coconut': 5,
    'papaya': 6,
    'orange': 7,
    'apple': 8,
    'muskmelon': 9,
    'watermelon': 10,
    'grapes': 11,
    'mango': 12,
    'banana': 13,
    'pomegranate': 14,
    'lentil': 15,
    'blackgram': 16,
    'mungbean': 17,
    'mothbeans': 18,
    'pigeonpeas': 19,
    'kidneybeans': 20,
    'chickpea': 21,
    'coffee': 22
}
crop['crop_num'] = crop['label'].map(crop_dict)
```

```
In [13]: crop['crop_num'].value_counts()
```

```
Out[13]: 1    100
         2    100
         3    100
         4    100
         5    100
         6    100
         7    100
         8    100
         9    100
        10    100
        11    100
        12    100
        13    100
        14    100
        15    100
        16    100
        17    100
        18    100
        19    100
        20    100
        21    100
        22    100
Name: crop_num, dtype: int64
```

```
In [14]: # crop.drop(['label'],axis=1,inplace=True)
         crop.head()
```

```
Out[14]:
```

	N	P	K	temperature	humidity	ph	rainfall	label	crop_num
0	90	42	43	20.879744	82.002744	6.502985	202.935536	rice	1
1	85	58	41	21.770462	80.319644	7.038096	226.655537	rice	1
2	60	55	44	23.004459	82.320763	7.840207	263.964248	rice	1
3	74	35	40	26.491096	80.158363	6.980401	242.864034	rice	1
4	78	42	42	20.130175	81.604873	7.628473	262.717340	rice	1

```
In [15]: #TRAIN, TEST, SPLIT
```

```
In [16]: X = crop.drop(['crop_num','label'],axis=1)
         y = crop['crop_num']
```

```
In [17]: X
```

Out[17]:

	N	P	K	temperature	humidity	ph	rainfall
0	90	42	43	20.879744	82.002744	6.502985	202.935536
1	85	58	41	21.770462	80.319644	7.038096	226.655537
2	60	55	44	23.004459	82.320763	7.840207	263.964248
3	74	35	40	26.491096	80.158363	6.980401	242.864034
4	78	42	42	20.130175	81.604873	7.628473	262.717340
...
2195	107	34	32	26.774637	66.413269	6.780064	177.774507
2196	99	15	27	27.417112	56.636362	6.086922	127.924610
2197	118	33	30	24.131797	67.225123	6.362608	173.322839
2198	117	32	34	26.272418	52.127394	6.758793	127.175293
2199	104	18	30	23.603016	60.396475	6.779833	140.937041

2200 rows × 7 columns

In [18]: `y.shape`Out[18]: `(2200,)`In [19]: `from sklearn.model_selection import train_test_split`In [20]: `X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)`In [21]: `X_train.shape`Out[21]: `(1760, 7)`In [22]: `X_test.shape`Out[22]: `(440, 7)`In [23]: `X_train`

Out[23]:

	N	P	K	temperature	humidity	ph	rainfall
1656	17	16	14	16.396243	92.181519	6.625539	102.944161
752	37	79	19	27.543848	69.347863	7.143943	69.408782
892	7	73	25	27.521856	63.132153	7.288057	45.208411
1041	101	70	48	25.360592	75.031933	6.012697	116.553145
1179	0	17	30	35.474783	47.972305	6.279134	97.790725
...
1638	10	5	5	21.213070	91.353492	7.817846	112.983436
1095	108	94	47	27.359116	84.546250	6.387431	90.812505
1130	11	36	31	27.920633	51.779659	6.475449	100.258567
1294	11	124	204	13.429886	80.066340	6.361141	71.400430
860	32	78	22	23.970814	62.355576	7.007038	53.409060

1760 rows × 7 columns

In [24]: *#Scale the features using MinMaxScaler*

```
In [25]: from sklearn.preprocessing import MinMaxScaler
ms = MinMaxScaler()

X_train = ms.fit_transform(X_train)
X_test = ms.transform(X_test)
```

In [26]: X_train

Out[26]:

```
array([[0.12142857, 0.07857143, 0.045      , ..., 0.9089898 , 0.48532225,
        0.29685161],
       [0.26428571, 0.52857143, 0.07      , ..., 0.64257946, 0.56594073,
        0.17630752],
       [0.05      , 0.48571429, 0.1       , ..., 0.57005802, 0.58835229,
        0.08931844],
       ...,
       [0.07857143, 0.22142857, 0.13      , ..., 0.43760347, 0.46198144,
        0.28719815],
       [0.07857143, 0.85      , 0.995     , ..., 0.76763665, 0.44420505,
        0.18346657],
       [0.22857143, 0.52142857, 0.085     , ..., 0.56099735, 0.54465022,
        0.11879596]])
```

In [27]: *#Standardization*

```
In [28]: from sklearn.preprocessing import StandardScaler
sc = StandardScaler()

sc.fit(X_train)
X_train = sc.transform(X_train)
X_test = sc.transform(X_test)
```

In [29]: X_train

```
Out[29]: array([[ -9.03426596e-01, -1.12616170e+00, -6.68506601e-01, ...,
          9.36586183e-01,  1.93473784e-01,  5.14970176e-03],
        [-3.67051340e-01,  7.70358846e-01, -5.70589522e-01, ...,
        -1.00470485e-01,  8.63917548e-01, -6.05290566e-01],
        [-1.17161422e+00,  5.89737842e-01, -4.53089028e-01, ...,
        -3.82774991e-01,  1.05029771e+00, -1.04580687e+00],
        ...,
        [-1.06433917e+00, -5.24091685e-01, -3.35588533e-01, ...,
        -8.98381379e-01, -6.34357580e-04, -4.37358211e-02],
        [-1.06433917e+00,  2.12501638e+00,  3.05234239e+00, ...,
        3.86340190e-01, -1.48467347e-01, -5.69036842e-01],
        [-5.01145154e-01,  7.40255346e-01, -5.11839275e-01, ...,
        -4.18045489e-01,  6.86860180e-01, -8.96531475e-01]])
```

```
In [30]: from sklearn.linear_model import LogisticRegression
from sklearn.naive_bayes import GaussianNB
from sklearn.svm import SVC
from sklearn.neighbors import KNeighborsClassifier
from sklearn.tree import DecisionTreeClassifier
from sklearn.tree import ExtraTreeClassifier
from sklearn.ensemble import RandomForestClassifier
from sklearn.ensemble import BaggingClassifier
from sklearn.ensemble import GradientBoostingClassifier
from sklearn.ensemble import AdaBoostClassifier
from sklearn.metrics import accuracy_score

# create instances of all models
models = {
    'Logistic Regression': LogisticRegression(),
    'Naive Bayes': GaussianNB(),
    'Support Vector Machine': SVC(),
    'K-Nearest Neighbors': KNeighborsClassifier(),
    'Decision Tree': DecisionTreeClassifier(),
    'Random Forest': RandomForestClassifier(),
    'Bagging': BaggingClassifier(),
    'AdaBoost': AdaBoostClassifier(),
    'Gradient Boosting': GradientBoostingClassifier(),
    'Extra Trees': ExtraTreeClassifier(),
}

for name, md in models.items():
    md.fit(X_train,y_train)
    ypred = md.predict(X_test)

    print(f"{name} with accuracy : {accuracy_score(y_test,ypred)}")
```

```
Logistic Regression with accuracy : 0.9636363636363636
Naive Bayes with accuracy : 0.9954545454545455
Support Vector Machine with accuracy : 0.9681818181818181
K-Nearest Neighbors with accuracy : 0.9590909090909091
Decision Tree with accuracy : 0.9818181818181818
```

C:\Users\Utkarsh\anaconda3\lib\site-packages\sklearn\neighbors_classification.py: 228: FutureWarning: Unlike other reduction functions (e.g. `skew`, `kurtosis`), the default behavior of `mode` typically preserves the axis it acts along. In SciPy 1.11.0, this behavior will change: the default value of `keepdims` will become False, the `axis` over which the statistic is taken will be eliminated, and the value None will no longer be accepted. Set `keepdims` to True or False to avoid this warning.

```
mode, _ = stats.mode(_y[neigh_ind, k], axis=1)
```


Random Forest with accuracy : 0.9954545454545455
 Bagging with accuracy : 0.9886363636363636
 AdaBoost with accuracy : 0.1409090909090909
 Gradient Boosting with accuracy : 0.9818181818181818
 Extra Trees with accuracy : 0.8818181818181818

```
In [31]: rfc = RandomForestClassifier()
         rfc.fit(X_train,y_train)
         ypred = rfc.predict(X_test)
         accuracy_score(y_test,ypred)
```

```
Out[31]: 0.9931818181818182
```

```
In [32]: #Predictive System
```

```
In [33]: def recommendation(N,P,k,temperature,humidity,ph,rainfal):
         features = np.array([[N,P,k,temperature,humidity,ph,rainfal]])
         transformed_features = ms.fit_transform(features)
         transformed_features = sc.fit_transform(transformed_features)
         prediction = rfc.predict(transformed_features).reshape(1,-1)

         return prediction[0]
```

```
In [ ]: N = 60
         P = 90
         k = 87
         temperature = 16.5
         humidity = 4.5
         ph = 30.0
         rainfall = 10.0

         predict = recommendation(N,P,k,temperature,humidity,ph,rainfall)

         crop_dict = {1: "Rice", 2: "Maize", 3: "Jute", 4: "Cotton", 5: "Coconut", 6: "Papaya",
                       8: "Apple", 9: "Muskmelon", 10: "Watermelon", 11: "Grapes", 12: "Mango",
                       14: "Pomegranate", 15: "Lentil", 16: "Blackgram", 17: "Mungbean",
                       19: "Pigeonpeas", 20: "Kidneybeans", 21: "Chickpea", 22: "Coffee"}

         if predict[0] in crop_dict:
             crop = crop_dict[predict[0]]
             print("{} is a best crop to be cultivated ".format(crop))
         else:
             print("Sorry are not able to recommend a proper crop for this environment")
```

```
In [ ]: import pickle
         pickle.dump(rfc,open('model.pkl','wb'))
         pickle.dump(ms,open('minmaxscaler.pkl','wb'))
         pickle.dump(sc,open('standscaler.pkl','wb'))
```

```
In [ ]:
```